

CSC/CPE 138

COMPUTER NETWORK FUNDAMENTALS




# Lecture 4\_1: The Network Layer (Data Plane)

California State University, Sacramento  
Fall 2024

Slide Courtesy: Computer Networking: A Top-Down Approach, Kurose Ross, 8<sup>th</sup> Edition

# Lecture 4\_1 Overview

- Network layer
  - data plane
  - control plane
- What's inside a router
  - input ports, switching, output ports
  - buffer management, scheduling
- IP: the Internet Protocol
  - Datagram format
  - IP addressing



# Two key network-layer functions

## Network-layer functions:

- *forwarding*: move packets from a router's input link to appropriate router output link
- *routing*: determine route taken by packets from source to destination
  - *routing algorithms*

## Analogy: taking a trip

- *forwarding*: process of getting through single interchange
- *routing*: process of planning trip from source to destination



forwarding




routing

# Network layer: data plane, control plane

## Data plane:

- *local*, per-router function
- determines how datagram arriving on router input port is forwarded to router output port

values in arriving  
packet header

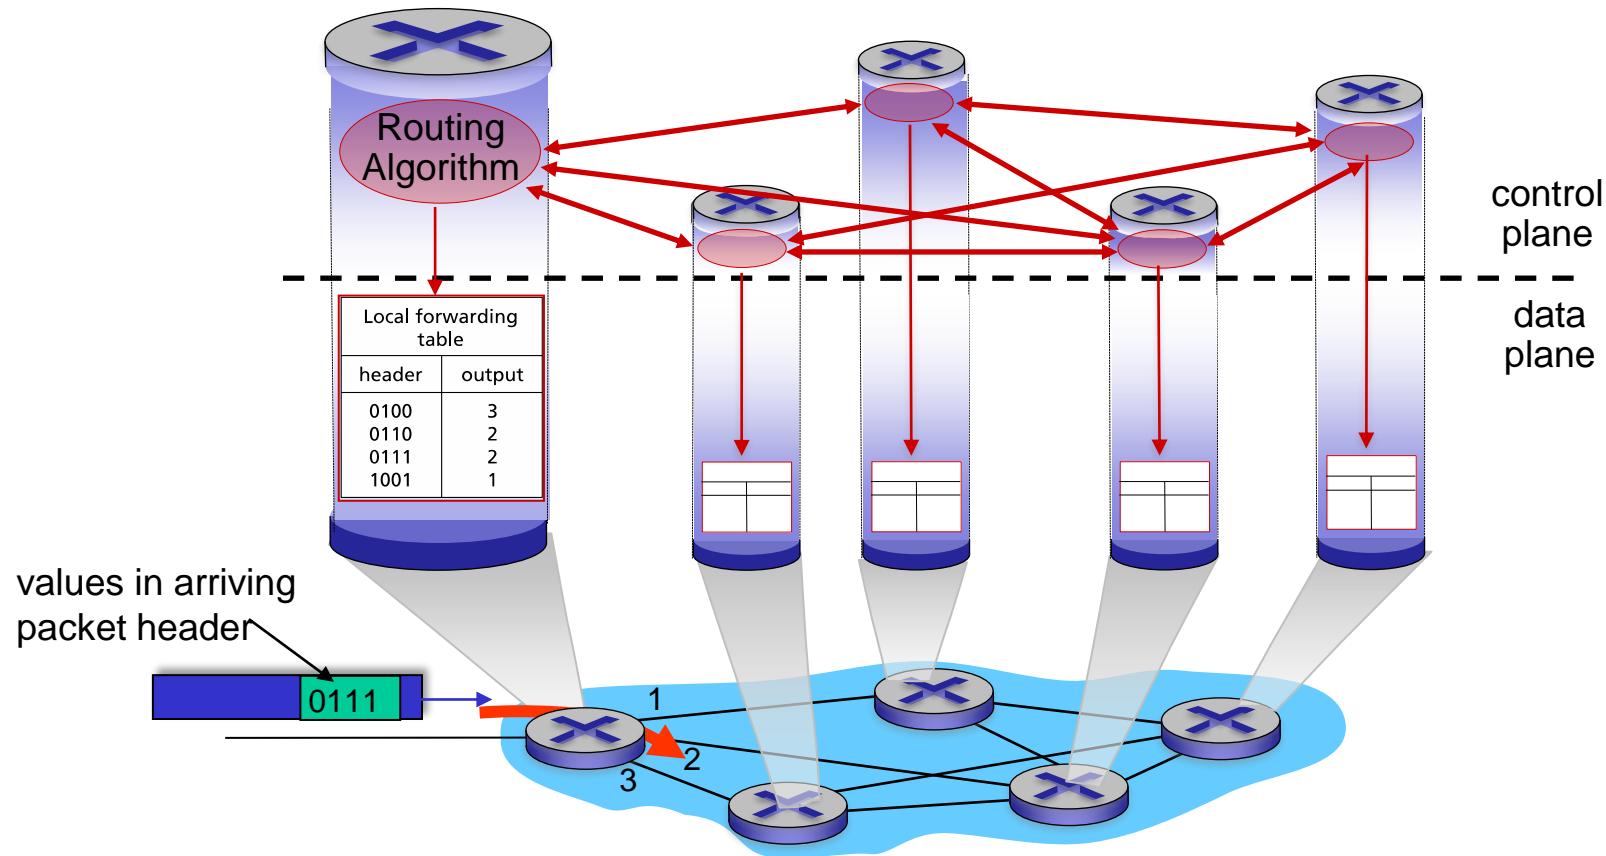


## Control plane

- *network-wide* logic
- determines how datagram is routed among routers along end-end path from source host to destination host
- two control-plane approaches:
  - *traditional routing algorithms*: implemented in routers
  - *software-defined networking (SDN)*: implemented in (remote) servers

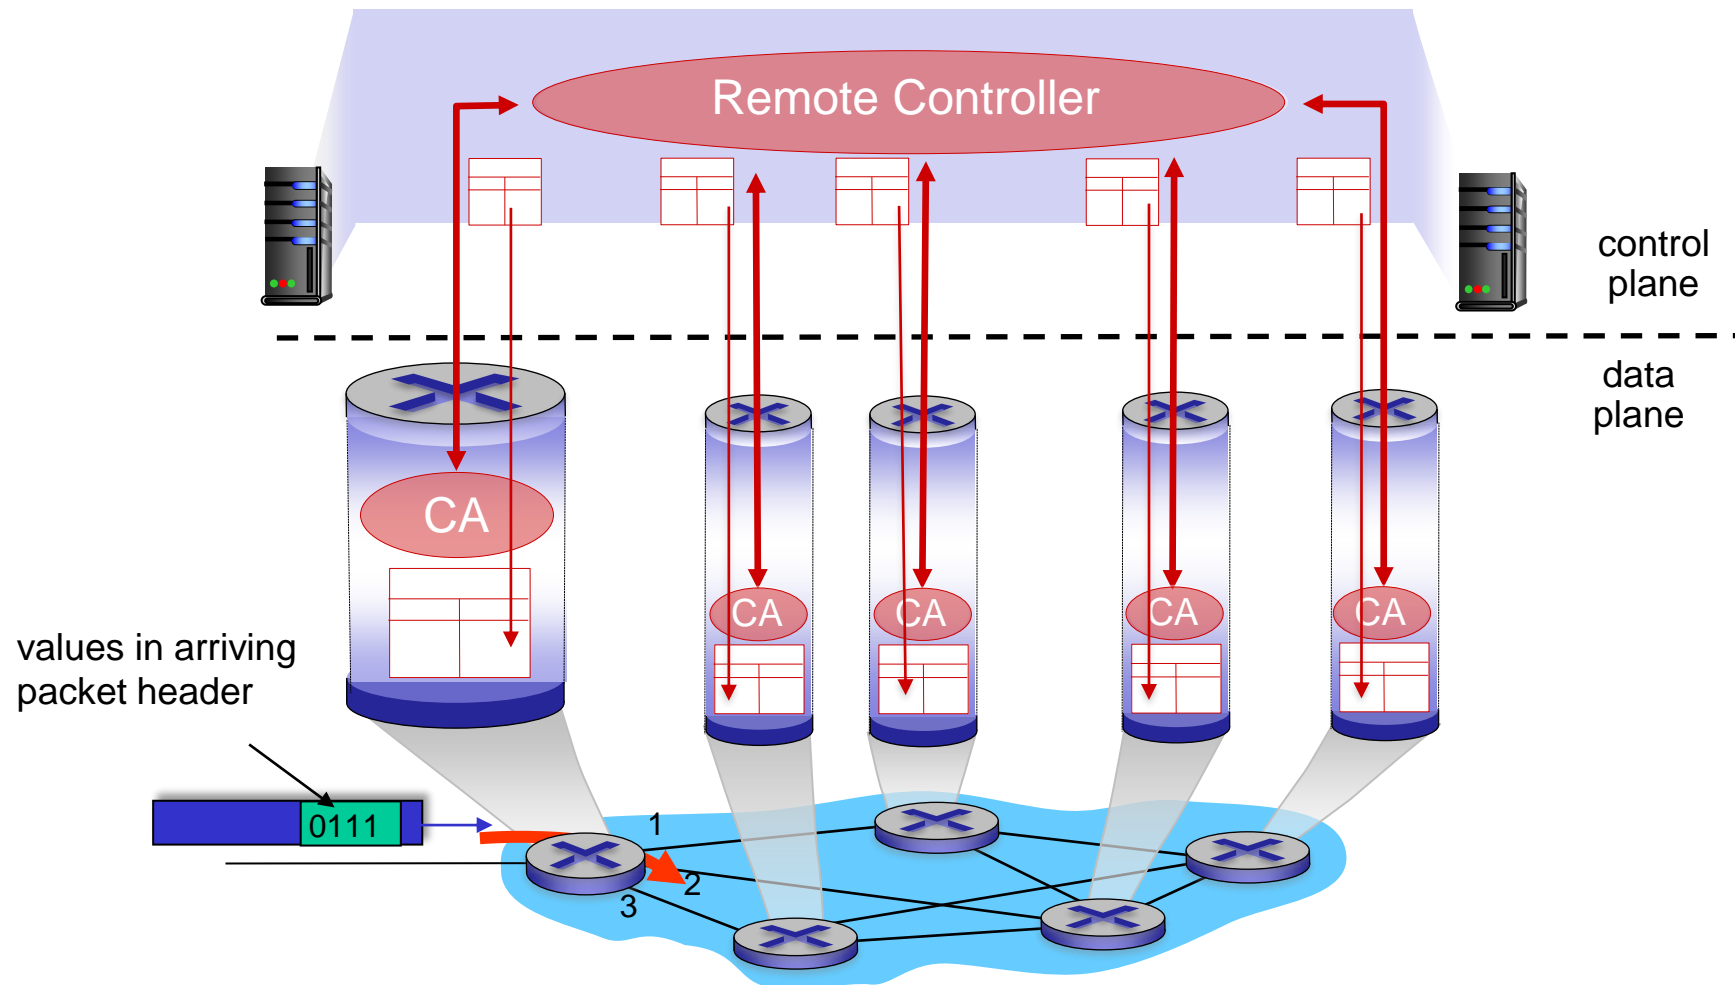
# Per-router control plane

Individual routing algorithm components *in each and every router* interact in the control plane



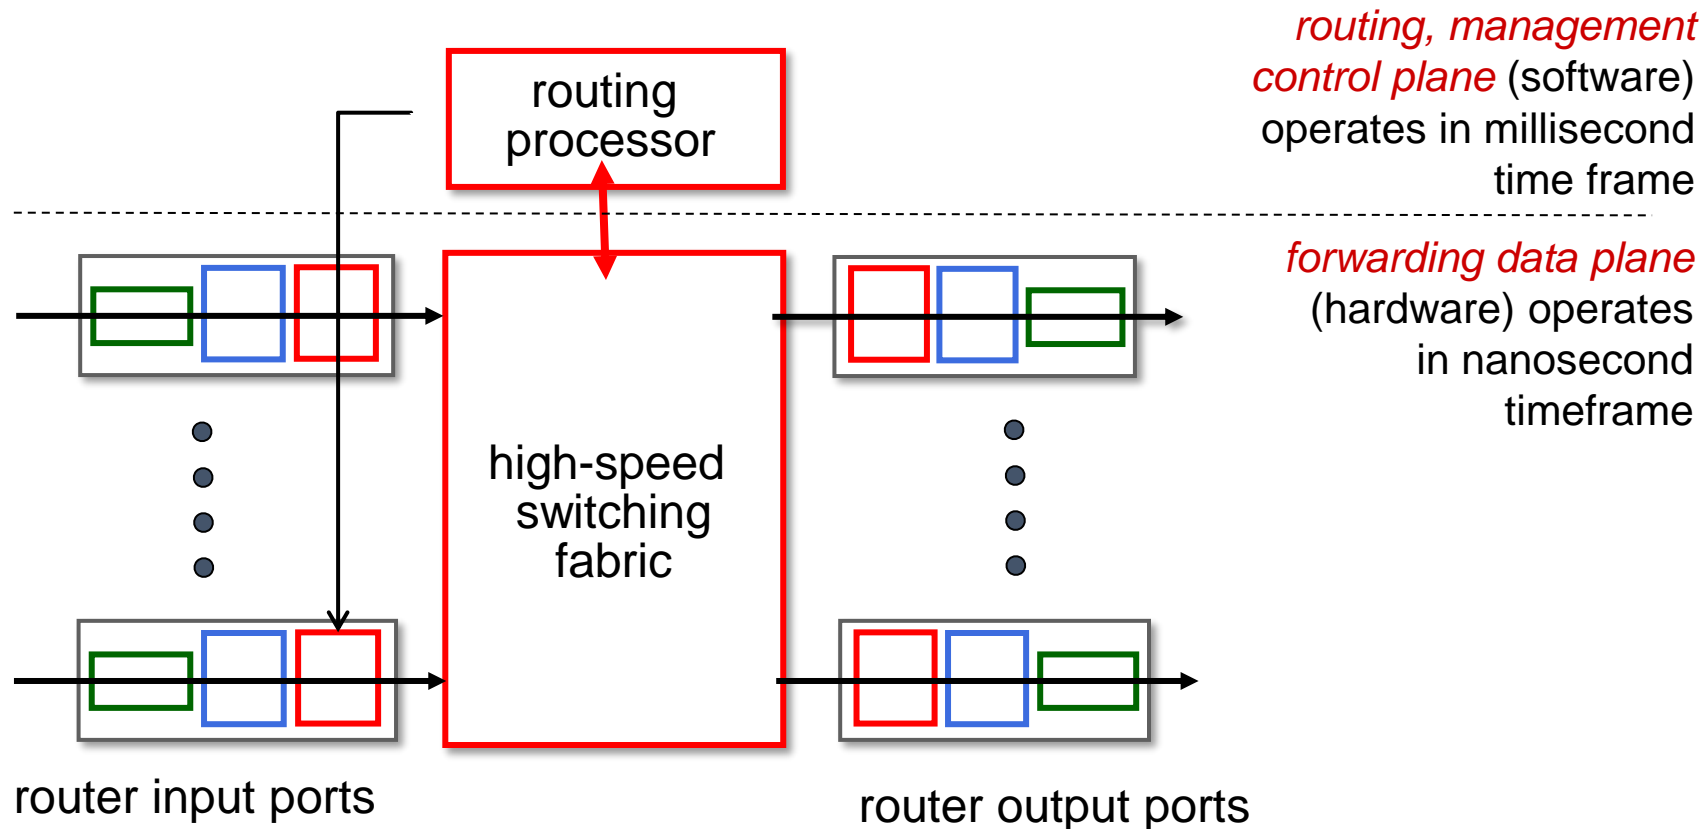
# Software-Defined Networking (SDN) Control Plane

Remote controller computes, installs forwarding tables in routers



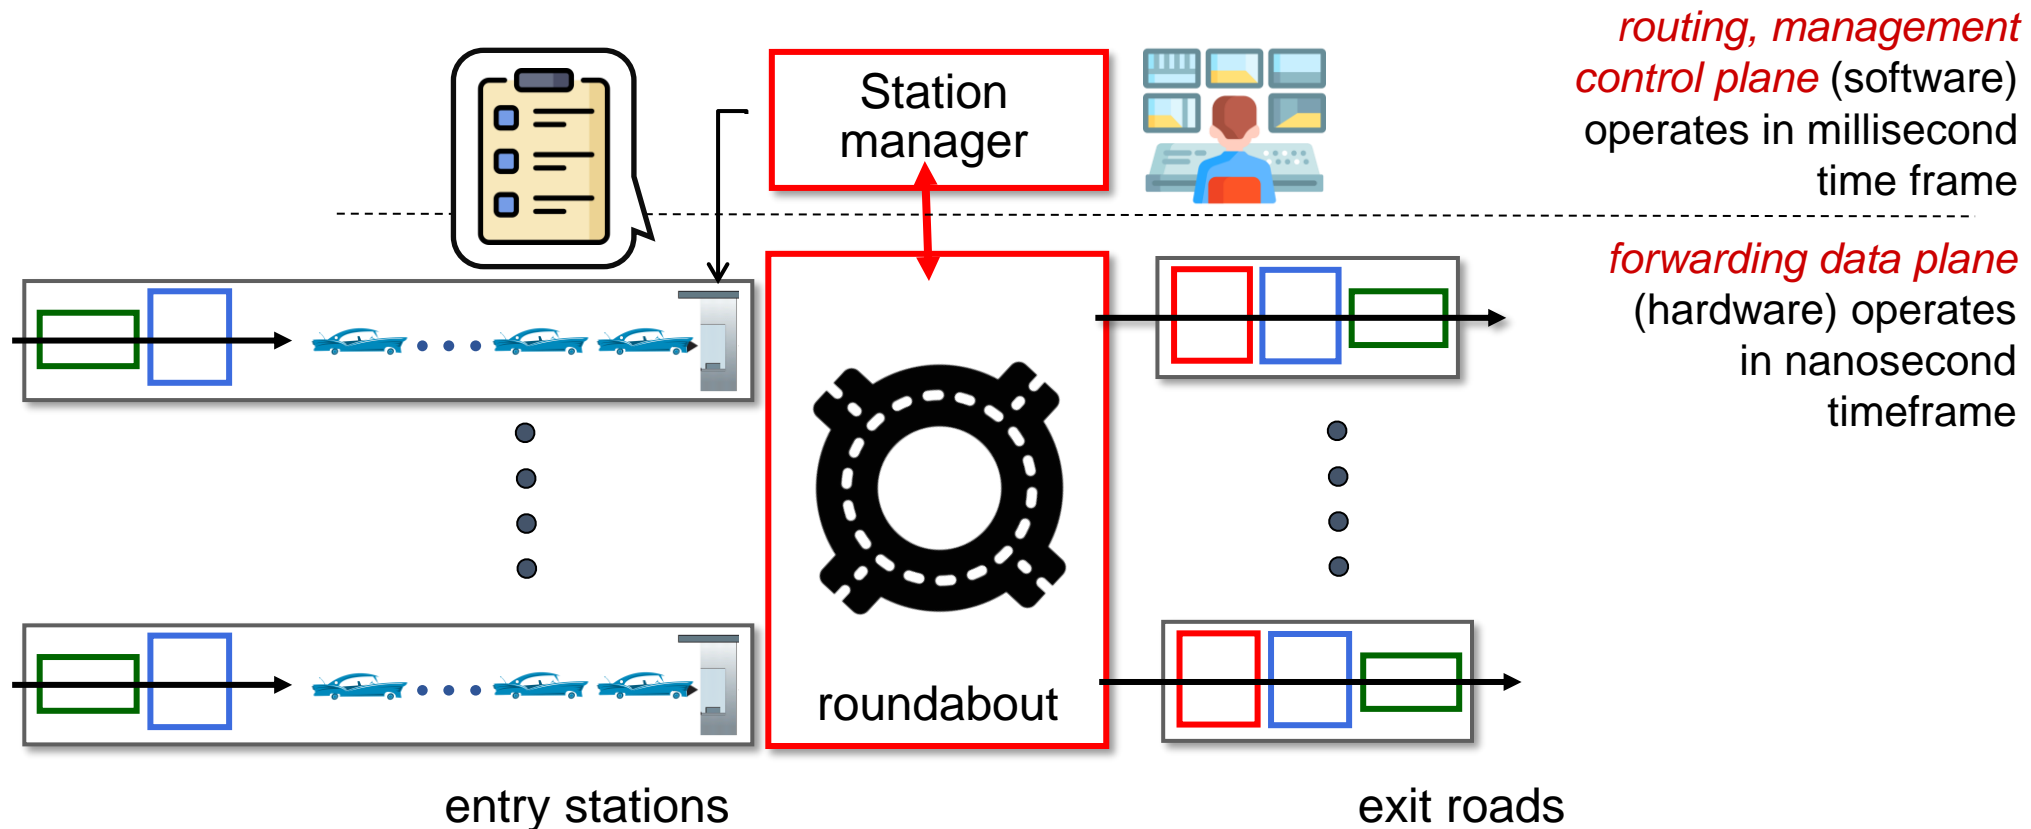
# Router architecture overview

high-level view of generic router architecture:



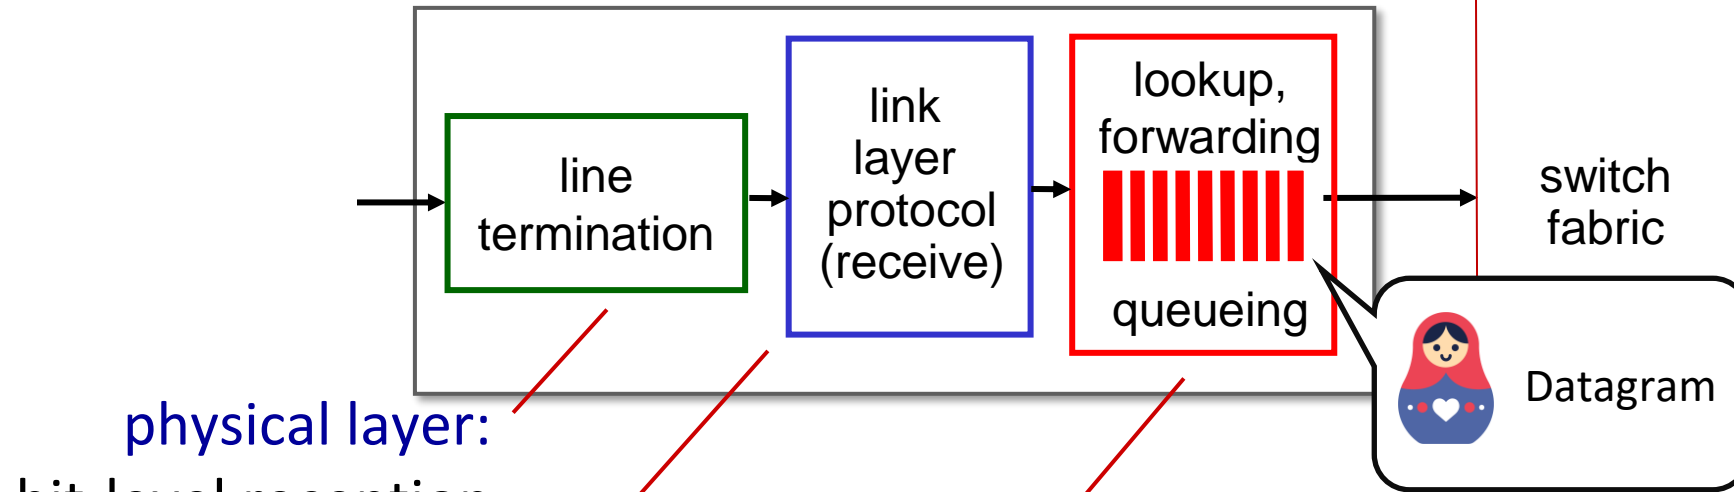
# Router architecture overview

analogy view of generic router architecture:





# Input port functions



physical layer:  
bit-level reception

link layer:  
e.g., Ethernet  
(chapter 6)

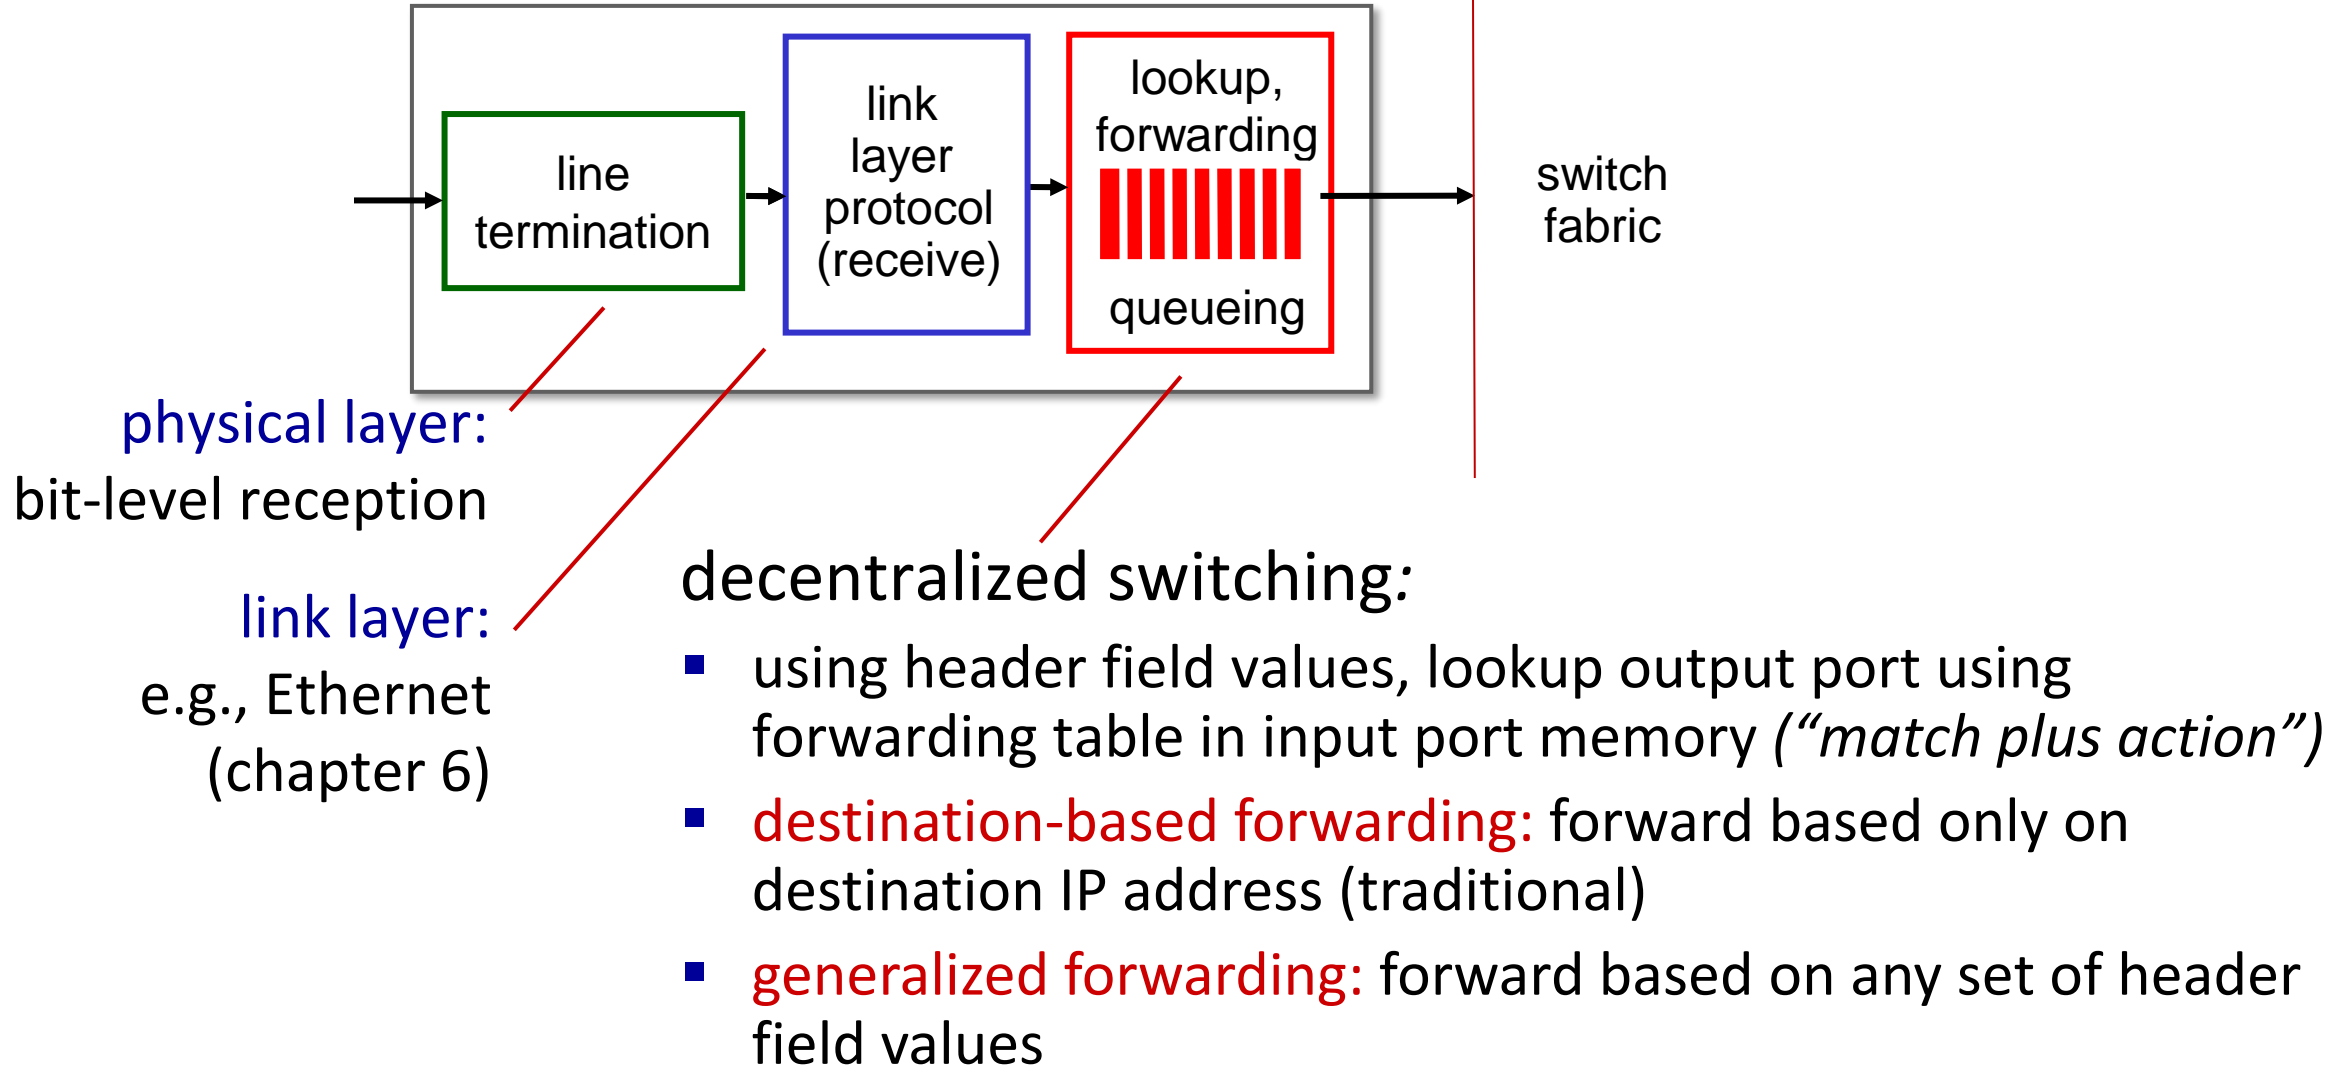


Frame

decentralized switching:

- using header field values, lookup output port using forwarding table in input port memory (*"match plus action"*)
- goal: complete input port processing at 'line speed'
- **input port queueing:** if datagrams arrive faster than forwarding rate into switch fabric

# Input port functions



# Destination-based forwarding

| <i>forwarding table</i>   |                |
|---|----------------|
| Destination Address Range   | Link Interface |
| 11001000 00010111 00010000 00000000<br>through<br>11001000 00010111 00010000 00000100 | 0              |
| 11001000 00010111 00010000 00000111<br>through<br>11001000 00010111 00011000 11111111 | 3              |
| 11001000 00010111 00011001 00000000<br>through<br>11001000 00010111 00011111 11111111 | 1              |
| 11001000 00010111 00011001 00000000<br>through<br>11001000 00010111 00011111 11111111 | 2              |
| otherwise   | 3              |

⚡ but what happens if ranges don't divide up so nicely:

# Longest prefix matching

## longest prefix match

when looking for forwarding table entry for given destination address, use *longest* address prefix that matches destination address.

| Destination Address Range                 | Link interface |
|---|----------------|
| 11001000    00010111    00010***    ***** | 0              |
| 11001000    00010111    00011000    ***** | 1              |
| 11001000    00010111    00011***    ***** | 2              |
| otherwise                                 | 3              |

examples:

11001000    00010111    00010110    10100001    which interface?

11001000    00010111    00011000    10101010    which interface?

# Longest prefix matching

## longest prefix match

when looking for forwarding table entry for given destination address, use *longest* address prefix that matches destination address.

| Destination Address Range               | Link interface |
|---|----------------|
| 11001000 00010111 00010*** *****        | 0              |
| 11001000 00010111 00011000 *****        | 1              |
| 11001000 <b>match!</b> 1 00011*** ***** | 2              |
| otherwise                               | 3              |

examples:

|                                     |                  |
|-------------------------------------|------------------|
| 11001000 00010111 00010110 10100001 | which interface? |
| 11001000 00010111 00011000 10101010 | which interface? |

# Longest prefix matching

## longest prefix match

when looking for forwarding table entry for given destination address, use *longest* address prefix that matches destination address.

| Destination Address Range        | Link interface |
|----------------------------------|----------------|
| 11001000 00010111 00010*** ***** | 0              |
| 11001000 00010111 00011000 ***** | 1              |
| 11001000 00010111 00011*** ***** | 2              |
| otherwise                        | 3              |

match!

examples:

|                                     |                  |
|-------------------------------------|------------------|
| 11001000 00010111 00010110 10100001 | which interface? |
| 11001000 00010111 00011000 10101010 | which interface? |

# Longest prefix matching

## longest prefix match

when looking for forwarding table entry for given destination address, use *longest* address prefix that matches destination address.

| Destination Address Range        | Link interface |
|----------------------------------|----------------|
| 11001000 00010111 00010*** ***** | 0              |
| 11001000 00010111 00011000 ***** | 1              |
| 11001000 00010111 00011*** ***** | 2              |
| otherwise                        | 3              |

match!

examples:

|                                     |                  |
|-------------------------------------|------------------|
| 11001000 00010111 00010110 10100001 | which interface? |
| 11001000 00010111 00011000 10101010 | which interface? |

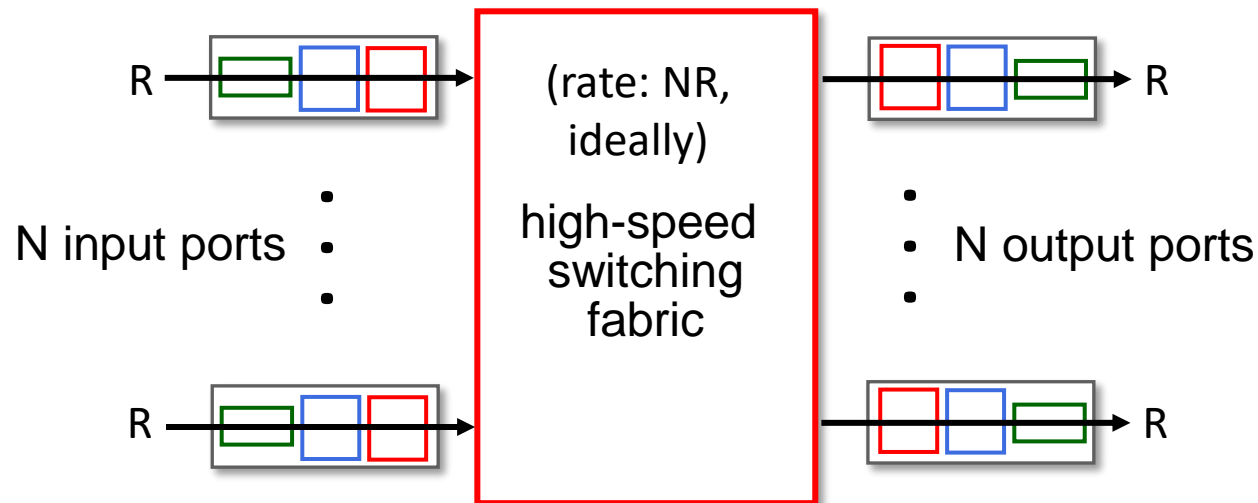
# Activity 4.1

- You are provided with below networks
  - 232.1.30.0
  - 232.1.35.0
  - 232.1.40.0
  - 232.1.42.0
  - 232.1.46.0
- Find the longest prefix match for 232.1.43.2






# Switching fabrics

- Transfer packet from input link to appropriate output link
- **Switching rate:** rate at which packets can be transfer from inputs to outputs
  - often measured as multiple of input/output line rate
  - N inputs: switching rate N times line rate desirable

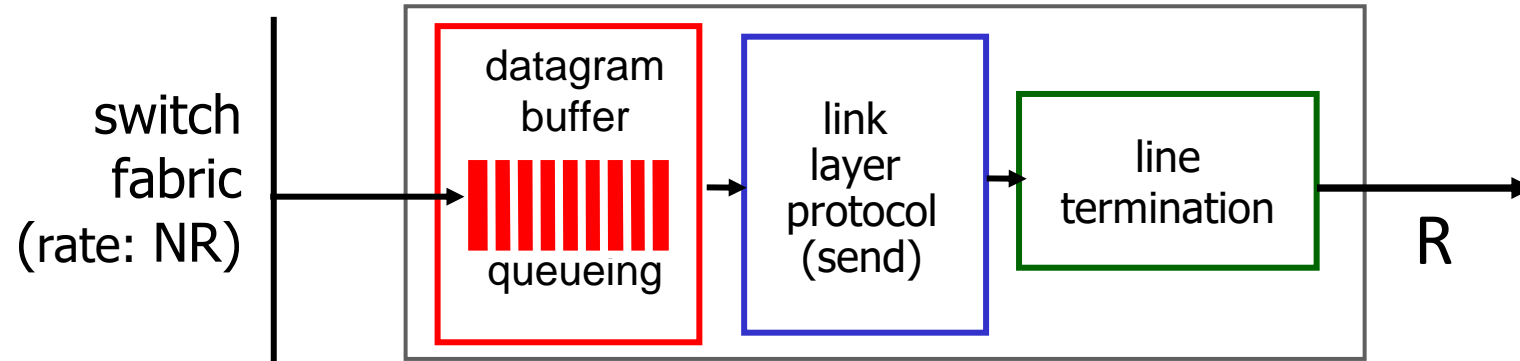


# Switching fabrics

- Transfer packet from input link to appropriate output link
- **Switching rate:** rate at which packets can be transfer from inputs to outputs
  - often measured as multiple of input/output line rate
  - N inputs: switching rate N times line rate desirable
- three major types of switching fabrics:



# Output port queuing



This is a really important slide

- **Buffering** required when datagrams arrive from fabric faster than link transmission rate. **Drop policy**: which datagrams to drop if no free buffers?
- **Scheduling discipline** chooses among queued datagrams for transmission

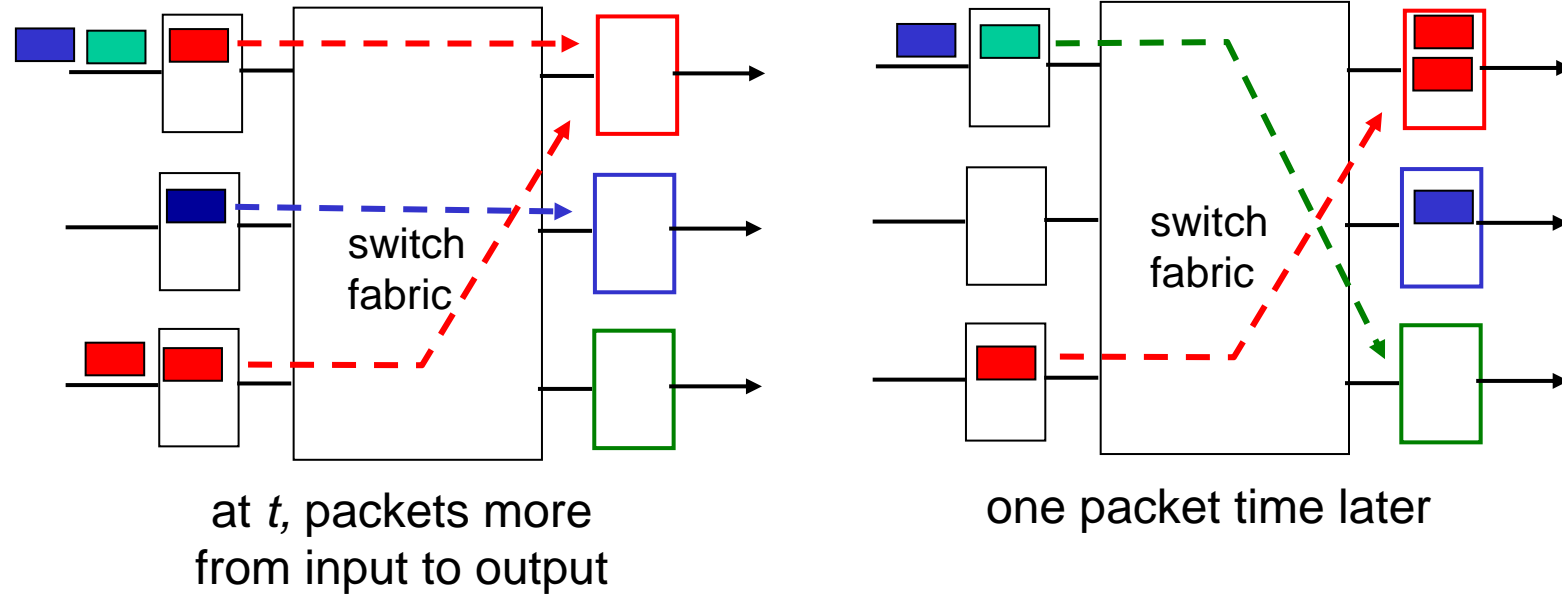


Datagrams can be lost due to congestion, lack of buffers



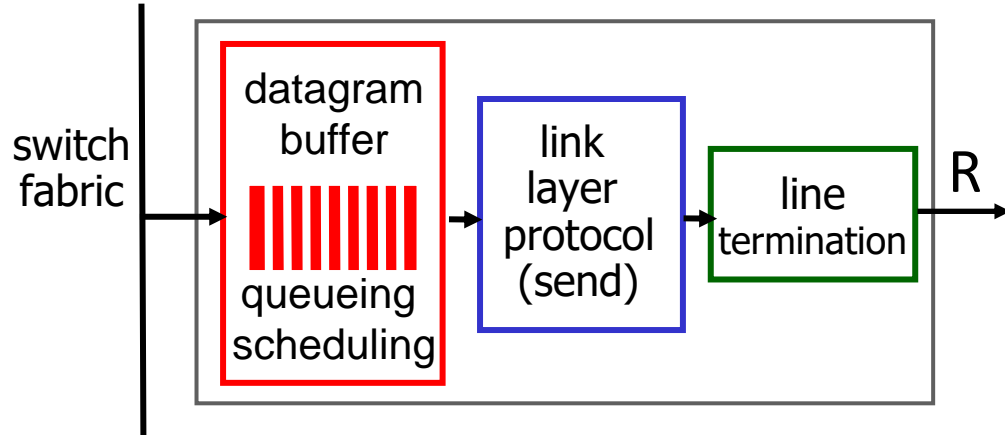
Priority scheduling – who gets best performance, network neutrality

# Output port queuing




- buffering when arrival rate via switch exceeds output line speed
- *queueing (delay) and loss due to output port buffer overflow!*

# Buffer Management



## Abstraction: queue



## buffer management:

- **drop**: which packet to add, drop when buffers are full
  - **tail drop**: drop arriving packet
  - **priority**: drop/remove on priority basis
- **marking**: which packets to mark to signal congestion (ECN, RED)

# Packet Scheduling: FCFS


**packet scheduling:** deciding which packet to send next on link

- first come, first served
- priority
- round robin
- weighted fair queueing

**FCFS:** packets transmitted in order of arrival to output port

- also known as: First-in-first-out (FIFO)
- real world examples?



Abstraction: queue



# Scheduling policies: priority

## *Priority scheduling:*


- arriving traffic classified, queued by class
  - any header fields can be used for classification
- send packet from highest priority queue that has buffered packets
  - FCFS within priority class



# Scheduling policies: round robin

## *Round Robin (RR) scheduling:*

- arriving traffic classified, queued by class
  - any header fields can be used for classification
- server cyclically, repeatedly scans class queues, sending one complete packet from each class (if available) in turn






# Scheduling policies: weighted fair queueing

## *Weighted Fair Queueing (WFQ):*

- generalized Round Robin
- each class,  $i$ , has weight,  $w_i$ , and gets weighted amount of service in each cycle:

$$\frac{w_i}{\sum_j w_j}$$

- minimum bandwidth guarantee (per-traffic-class)



## Activity 4.2

- Scheduling policies for packets
  - Each packet has a processing delay of 2 seconds
  - Packets 1,2,3,4 and 5 arriving at 0<sup>th</sup> ,1<sup>st</sup> ,2<sup>nd</sup> ,7<sup>th</sup> and 12<sup>th</sup> seconds with priority 10,40,20,40,20
  - Use FCFS and Priority scheduling to process packets

# Activity 4.2 Continued

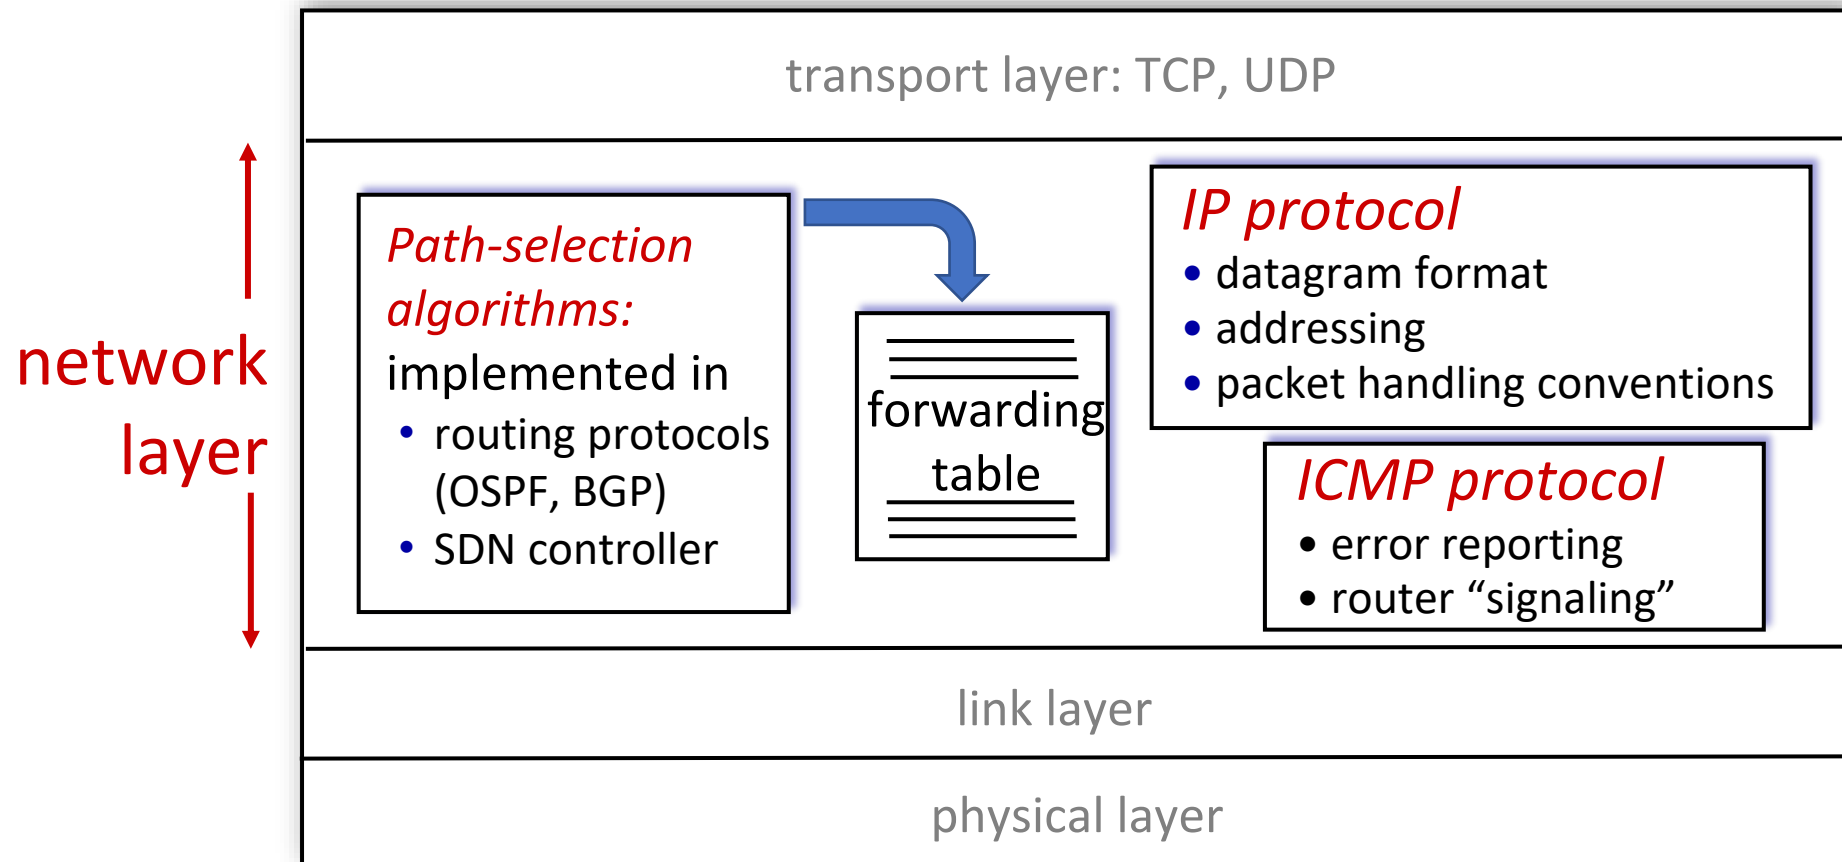
- Use round robin and WFQ scheduling for below packets to be processed

| Class | Weight | Packets |
|-------|--------|---------|
| 3     | 6      | 24      |
| 2     | 3      | 15      |
| 1     | 1      | 5       |

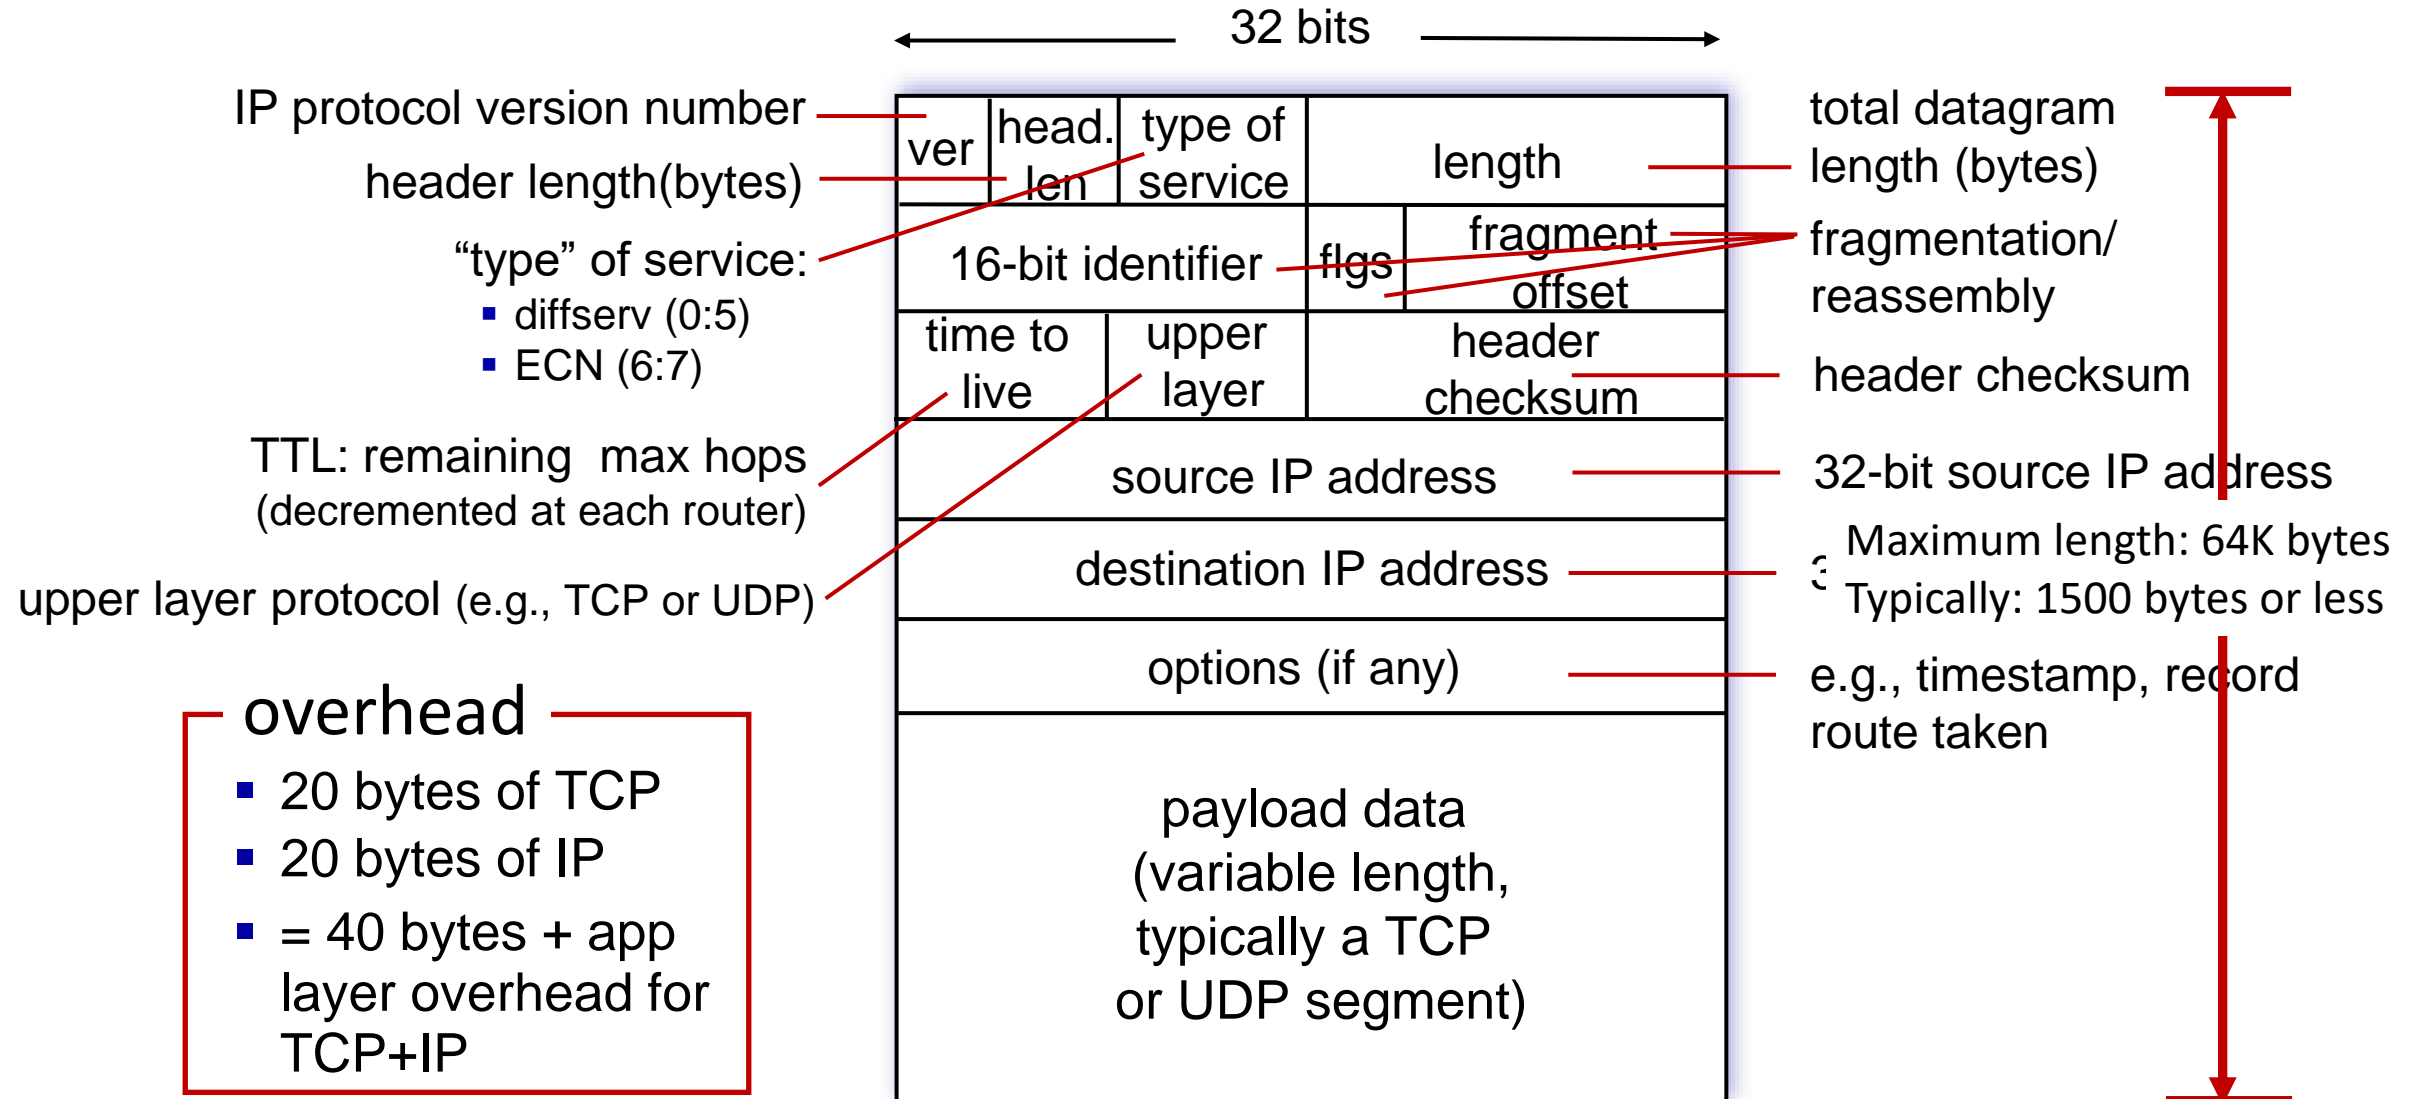
- For round robin, consider 1 packet from each class being processed as it does not have any weight
- Draw and time graph and identify the time in seconds for completion of above scheduling for round robin and WFQ policy if each packet takes 1 second to complete processing.

# Network Layer: Internet

host, router network layer functions:




# IP Datagram format (Read Textbook)



# IP addressing: introduction

- **IP address:** 32-bit identifier associated with each host or router *interface*
- **interface:** connection between host/router and physical link
  - router's typically have multiple interfaces
  - host typically has one or two interfaces (e.g., wired Ethernet, wireless 802.11)




dotted-decimal IP address notation:

223.1.1.1 =  $\underbrace{11011111}_{223} \underbrace{00000001}_1 \underbrace{00000001}_1 \underbrace{00000001}_1$

# IP addressing: introduction

- **IP address:** 32-bit identifier associated with each host or router *interface*
- **interface:** connection between host/router and physical link
  - router's typically have multiple interfaces
  - host typically has one or two interfaces (e.g., wired Ethernet, wireless 802.11)



dotted-decimal IP address notation:

223.1.1.1 = 11011111 00000001 00000001 00000001

223                      1                      1                      1


# IP addressing: introduction

**Q:** how are interfaces  
actually connected?

**A:** we'll learn about  
that in chapters 6, 7

*For now:* don't need to worry  
about how one interface is  
connected to another (with no  
intervening router)

**A:** wired  
Ethernet interfaces  
connected by  
Ethernet switches




**A:** wireless WiFi interfaces  
connected by WiFi base station



# Subnets

- *What's a subnet ?*
  - device interfaces that can physically reach each other **without passing through an intervening router**
- IP addresses have structure:
  - **subnet part**: devices in same subnet have common high order bits
  - **host part**: **remaining** low order bits




network consisting of 3 subnets



# Subnets

## *Recipe for defining subnets:*

- detach each interface from its host or router, creating “islands” of isolated networks
- each isolated network is called a *subnet*




subnet mask: /24

(high-order 24 bits: subnet part of IP address)

# Subnets


- where are the subnets?
- what are the /24 subnet addresses?



# IP addressing: CIDR

**CIDR: Classless InterDomain Routing** (pronounced “cider”)

- subnet portion of address of arbitrary length
- address format: **a.b.c.d/x**, where x is # bits in subnet portion of address



# Activity 4.3

## ■ IP Subnetting

- Given an IP address of 192.168.1.0/24
  - Determine the network address, subnet mask and usable hosts
  - Create two networks out of 192.168.1.0 of 110 hosts each and identify the network addresses, subnet masks and usable hosts.

# Lecture 4\_1 Summary

- Principles behind network layer services
  - Network layer service models
  - Forwarding versus routing
  - How a router works
  - IP Addressing and Subnetting